

Great Minds · Product Overview

A research tool and thought partner for people who work with ideas.

What It Is

Great Minds is a personal and team knowledge base powered by AI. You feed it your documents · articles, books, essays, notes, news clippings, URLs · and it reads them, organizes them, and becomes a research partner that can answer questions grounded in *your* materials, not the open internet.

Think of it as a research assistant that has actually read everything in your library and can pull connections across sources on demand.

How It Works

1. Build Your Brain

A "brain" is your knowledge base. You create one and start adding sources:

- **Drop files** · PDFs, text documents, essays, book excerpts
- **Paste URLs** · web articles are automatically converted and stored
- **Write directly** · add your own notes, ideas, or analysis

Once ingested, Great Minds automatically generates a **wiki** · a set of interconnected summary articles that map the terrain of your materials. Documents link to related documents. Concepts cross-reference each other. The tool builds structure from your raw inputs.

2. Ask Questions

The primary interface is a search bar: *"Ask a question across the knowledge base."*

You type a research question, and Great Minds searches your materials, reads the relevant sources, and synthesizes an answer with citations back to specific documents.

Every answer shows its work. You can see which documents were consulted, which searches were performed, and click through to read the original source material. Nothing is a black box.

3. How the AI Finds What It Needs

Under the hood, the AI doesn't just pattern-match against your question. It *reasons* about where to look, then actively navigates your knowledge base using three specialized retrieval tools:

- **search_wiki** · Hybrid search combining keyword matching (BM25) and semantic vector search across all wiki articles. The AI uses this when it needs to find which articles cover a topic, returning matching excerpts with article paths and section headings. This means it finds relevant material even when your question uses different vocabulary than your sources.
- **read_document** · Reads any document in the knowledge base by path · wiki articles, raw texts, notes, news clippings. When the AI reads a wiki article, it also receives that article's **forward links** (what it cites) and **backlinks** (what cites it), giving it a live view of the knowledge graph. It can follow these links to traverse from a summary article into the primary source material it was built from, or discover related articles it hadn't considered. Documents over 20KB are truncated with a prompt to request specific sections, so the AI reads selectively rather than dumping entire books into context.
- **query_documents** · Structured metadata search across the full document corpus. The AI can filter by author, publication date, genre (theoretical, polemical, etc.), content type (texts, news, ideas),

tags, and concepts. The available vocabulary for tags and concepts is injected dynamically from the database at query time, so the AI always knows exactly what's in the collection. This is how it answers questions like "what did [author] write about [topic] before [year]" · not by reading everything, but by narrowing the field first.

These tools run in a loop. The AI might start with a wiki search, read two articles, follow a backlink to a third, then query for all documents by a specific author to fill in the picture · making as many passes as it needs before composing an answer. Every tool call is logged, so you can inspect the full retrieval chain in the "thinking" section of any response.

4. Go Deeper

Two features support the natural rhythm of research · where one answer raises new questions:

- **Follow-ups** · Highlight a passage in an answer, add context, and ask a targeted follow-up question within the same session. The AI retains the full thread of your inquiry.
- **"By The Way" (BTW)** · A lightweight side-thread for quick clarifications. See a term you don't recognize or a claim you want unpacked? Highlight it, click BTW, and get a short explanation without leaving your main line of research. These are designed for 2-3 paragraph responses, not full investigations.

5. Browse and Read

Beyond Q&A, you can browse your knowledge base directly:

- **Wiki view** · Navigate the auto-generated wiki articles, follow links between topics, and see how your materials connect
 - **Document reader** · Read any source document in full, with backlinks showing which other articles reference it
 - **Session history** · Every research session is saved and searchable, so you can return to past lines of inquiry
-

Team Use

Great Minds supports shared brains with role-based access:

- **Owners** manage the brain, approve new content, and control membership
- **Editors** add documents and submit proposals for new content
- **Viewers** have read-only access to query and browse

The **proposals** workflow lets team members submit new documents or wiki articles for review before they enter the shared knowledge base. This keeps the corpus curated while allowing distributed contribution · useful for organizations where multiple people are gathering and analyzing source material.

What Makes It Different

Grounded answers, not general knowledge. Most AI tools draw from their training data · the entire internet, roughly. Great Minds answers *only* from your documents. If your corpus doesn't cover a topic, it says so. This is the difference between a tool that sounds confident and a tool you can actually trust for research.

Graph-aware retrieval. The AI doesn't just search · it navigates. It follows links between wiki articles, traces citations back to primary sources, and discovers related material through backlinks. Your knowledge base isn't a flat

pile of documents; it's a graph, and the AI traverses it the way a researcher would.

Transparent sourcing. Every claim traces back to a document you provided. You can verify, challenge, or follow the citation chain · the same standard you'd hold a human research assistant to.

Structure emerges from your materials. The auto-generated wiki isn't a summary dump. It's an interconnected map of ideas, with links between related topics, that evolves as you add more sources. It surfaces connections you might not have seen.

Sessions as research artifacts. Your question-and-answer threads are preserved as durable records. A month later, you can search your past sessions the same way you'd search your documents · your research process becomes part of your knowledge base.

Current State

Great Minds is a working application with a web interface, user accounts, and a backend API. The core research loop · ingest sources, ask questions, get grounded answers, go deeper · is fully functional. Team brains with proposals and role-based access are built. The wiki auto-generation and document linting (detecting broken links, orphaned documents, missing citations) are operational.

Some capabilities are available through the API but don't yet have buttons in the web interface · brain management, manual compilation triggers, and the full proposal review workflow are among them. These are built and working, just waiting on UI.

Who It's For

Great Minds is built for people who accumulate knowledge with a purpose: organizers building strategic analysis from many sources, researchers synthesizing across a body of literature, policy teams maintaining institutional knowledge, or anyone who needs to think rigorously across a growing corpus of materials · and wants a tool that thinks *with* them rather than *for* them.